

# Mathematics for ML



Srinivasa. R  
Bangalore, India

# Notation

- $a, b, c$  Scalar (integer or real)
- $\mathbf{x}, \mathbf{y}, \mathbf{z}$  Vector (bold-font, lower case)
- $\mathbf{A}, \mathbf{B}, \mathbf{C}$  Matrix (bold-font, upper-case)
- $\mathbf{A}, \mathbf{B}, \mathbf{C}$  Tensor ((bold-font, upper-case)
- $X, Y, Z$  Random variable (normal font, upper-case)
- $a \in \mathcal{A}$  Set membership:  $a$  is member of set  $\mathcal{A}$
- $|\mathcal{A}|$  Cardinality: number of items in set  $\mathcal{A}$
- $\|\mathbf{w}\|$  Norm of vector  $\mathbf{w}$
- $\mathbf{u} \cdot \mathbf{w}$  or  $\langle \mathbf{u}, \mathbf{w} \rangle$  Dot product of vectors  $\mathbf{u}$  and  $\mathbf{w}$
- $\mathbb{R}$  Set of real numbers
- $\mathbb{R}^n$  Real numbers space of dimension  $n$
- $y = f(x)$  or  $x \mapsto f(x)$  Function (map): assign a unique value  $f(x)$  to each input value  $x$
- $f: \mathbb{R}^n \rightarrow \mathbb{R}$  Function (map): map an  $n$ -dimensional vector into a scalar

# Notation

- $\mathbf{A} \odot \mathbf{B}$  Element-wise product of matrices  $\mathbf{A}$  and  $\mathbf{B}$
- $\mathbf{A}$  Pseudo-inverse of matrix  $\mathbf{A}$
- $\frac{d^n f}{dx^n}$   $n$ -th derivative of function  $f$  with respect to  $x$
- $\nabla f(\mathbf{x})$  Gradient of function  $f$  with respect to  $\mathbf{x}$
- $\mathbf{H}_f$  Hessian matrix of function  $f$
- $X \sim P$  Random variable  $X$  has distribution  $P$
- $P(X|Y)$  Probability of  $X$  given  $Y$
- $\mathcal{N}(\mu, \sigma^2)$  Gaussian distribution with mean  $\mu$  and variance  $\sigma^2$
- $\mathbb{E}_{X \sim P}[f(X)]$  Expectation of  $f(X)$  with respect to  $P(X)$
- $\text{Var}(f(X))$  Variance of  $f(X)$
- $\text{Cov}(f(X), g(Y))$  Covariance of  $f(X)$  and  $g(Y)$
- $\text{corr}(X, Y)$  Correlation coefficient for  $X$  and  $Y$
- $D_{\text{KL}}(P \parallel Q)$  Kullback-Leibler divergence for distributions  $P$  and  $Q$
- $\mathcal{H}(P, Q)$  Cross-entropy for distributions  $P$  and  $Q$

# Vectors

- **Vector** definition
  - **Computer science:** *vector* is a one-dimensional array of ordered real-valued scalars
  - **Mathematics:** *vector* is a quantity possessing both magnitude and direction, represented by an arrow indicating the direction, and the length of which is proportional to the magnitude
- Vectors are written in column form or in row form
  - Denoted by bold-font lower-case letters




$$\mathbf{x} = \begin{bmatrix} 1 \\ 7 \\ 0 \\ 1 \end{bmatrix} \qquad \mathbf{x} = [1 \quad 7 \quad 0 \quad 1]^T$$

- For a general form vector with  $m$  elements, the vector lies in the  $m$ -dimensional space  $\mathbf{x} \in \mathbb{R}^m$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix}$$



# Geometry of Vectors

- First interpretation of a vector: **point in space**
  - E.g., in 2D we can visualize the data points with respect to a coordinate origin
- Second interpretation of a vector: **direction in space**
  - E.g., the vector  $\mathbf{w} = [3, 2]^T$  has a direction of 3 steps to the right and 2 steps up
  - The notation  $\mathbf{w}$  is sometimes used to indicate that the vectors have a direction
  - All vectors in the figure have the same direction
- Vector **addition**
  - We add the coordinates, and follow the directions given by the two vectors that are added



# Contd...


- The geometric interpretation of vectors as points in space allow us to consider a training set of input examples in ML as a **collection of points in space**
  - Hence, classification can be viewed as discovering how to separate two clusters of points belonging to different classes (left picture)
    - Rather than distinguishing images containing cars, planes, buildings, for example
  - Or, it can help to visualize zero-centering and normalization of training data (right picture)



# Dot Product and Angles

- **Dot product** of vectors,  $\mathbf{u} \cdot \mathbf{w} = \mathbf{u}^T \mathbf{w} = \sum_{ii} u_i w_i$ 
  - It is also referred to as **inner product**, or **scalar product** of vectors
  - The dot product  $\mathbf{u} \cdot \mathbf{w}$  is also often denoted by  $\langle \mathbf{u}, \mathbf{w} \rangle$
- The dot product is a symmetric operation,  $\mathbf{u} \cdot \mathbf{w} = \mathbf{u}^T \mathbf{w} = \mathbf{w}^T \mathbf{u} = \mathbf{w} \cdot \mathbf{u}$
- Geometric interpretation of a dot product:  
**angle** between two vectors
  - I.e., dot product  $\mathbf{u} \cdot \mathbf{w}$  over the norms of the vectors is  $\cos(\theta)$

$$\mathbf{u} \cdot \mathbf{w} = \|\mathbf{u}\| \|\mathbf{w}\| \cos(\theta) \quad \cos \theta = \frac{\mathbf{u} \cdot \mathbf{w}}{\|\mathbf{u}\| \|\mathbf{w}\|}$$



- If two vectors are orthogonal:  $\theta = 90^\circ$ , i.e.,  $\cos(\theta) = 0$ , then  $\mathbf{u} \cdot \mathbf{w} = 0$
- Also, in ML the term  $\cos \theta = \frac{\mathbf{u} \cdot \mathbf{w}}{\|\mathbf{u}\| \|\mathbf{w}\|}$  is sometimes employed as a measure of closeness of two vectors/data instances, and it is referred to as **cosine similarity**

# Norm of a Vector

## Vectors

- A vector **norm** is a function that maps a vector to a scalar value
  - The norm is a measure of the size of the vector
- The norm  $ff$  should satisfy the following properties:
  - Scaling:  $ff(\alpha \mathbf{x}) = |\alpha| ff(\mathbf{x})$
  - Triangle inequality:  $ff(\mathbf{x} + \mathbf{y}) \leq ff(\mathbf{x}) + ff(\mathbf{y})$
  - Must be non-negative:  $ff(\mathbf{x}) \geq 0$
- The general  $\ell_p$  norm of a vector  $\mathbf{x}$  is obtained as: 
$$\|\mathbf{x}\|_p = \left( \sum_{i=1}^m |x_i|^p \right)^{\frac{1}{p}}$$
  - On next page we will review the most common norms, obtained for  $p = 1, 2$ , and  $\infty$



# Contd...

## Vectors

- For  $p=2$ , we have  $\ell_2$  norm

- Also called **Euclidean norm**
- It is the most often used norm
- $\ell_2$  norm is often denoted just as  $\|\mathbf{x}\|$  with the subscript 2 omitted

$$\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^m x_i^2} = \sqrt{\mathbf{x}^T \mathbf{x}}$$

- For  $p=1$ , we have  $\ell_1$  norm

- Uses the absolute values of the elements
- Discriminate between zero and non-zero elements

$$\|\mathbf{x}\|_1 = \sum_{i=1}^m |x_i|$$

- For  $p=\infty$ , we have  $\ell_\infty$  norm

- Known as **infinity norm**, or **max norm**
- Outputs the absolute value of the largest element

$$\|\mathbf{x}\|_\infty = \max_i |x_i|$$

- $\ell_0$  norm outputs the number of non-zero elements


- It is not an  $\ell_p$  norm, and it is not really a norm function either (it is incorrectly called a norm)

# Vector Projection

## Vectors

- **Orthogonal projection** of a vector  $\mathbf{y}$  onto vector  $\mathbf{x}$ 
  - The projection can take place in any space of dimensionality  $\geq 2$
  - The **unit vector** in the direction of  $\mathbf{x}$  is  $\frac{\mathbf{x}}{\|\mathbf{x}\|}$ 
    - A unit vector has norm equal to 1
  - The length of the projection of  $\mathbf{y}$  onto  $\mathbf{x}$  is  $\|\mathbf{y}\| \cos(\theta)$
  - The orthogonal project is the vector  $\text{proj}_{\mathbf{x}}(\mathbf{y})$

$$\text{proj}_{\mathbf{x}}(\mathbf{y}) = \frac{\mathbf{x} \|\mathbf{y}\| \cos(\theta)}{\|\mathbf{x}\|}$$




# Hyperplanes


## Hyperplanes

- **Hyperplane** is a subspace whose dimension is one less than that of its ambient space
  - In a 2D space, a hyperplane is a straight line (i.e., 1D)
  - In a 3D, a hyperplane is a plane (i.e., 2D)
  - In a  $d$ -dimensional vector space, a hyperplane has  $d-1$  dimensions, and divides the space into two half-spaces
- Hyperplane is a generalization of a concept of plane in high-dimensional space
- In ML, hyperplanes are **decision boundaries** used for linear classification
  - Data points falling on either sides of the hyperplane are attributed to different classes

A hyperplane in  $\mathbb{R}^2$  is a line



A hyperplane in  $\mathbb{R}^3$  is a plane



# Contd...

## Hyperplanes

- For example, for a given data point  $\mathbf{w} = [2, 1]^T$ , we can use dot-product to find the hyperplane for which  $\mathbf{v} \cdot \mathbf{w} = 1$

- I.e., all vectors with  $\mathbf{v} \cdot \mathbf{w} > 1$  can be classified as one class, and all vectors with  $\mathbf{v} \cdot \mathbf{w} < 1$  can be classified as another class



- Solving  $\mathbf{v} \cdot \mathbf{w} = 1$ , we obtain

$$\|\mathbf{v}\| \|\mathbf{w}\| \cos(\theta) = 1 \iff \|\mathbf{v}\| \cos(\theta) = \frac{1}{\|\mathbf{w}\|} = \frac{1}{\sqrt{5}}$$

- I.e., the solution is the set of points for which  $\mathbf{v} \cdot \mathbf{w} = 1$  meaning the points lay on the line that is orthogonal to the vector  $\mathbf{w}$

- That is the line  $2x + y = 1$


- The orthogonal projection of  $\mathbf{w}$  onto  $\mathbf{v}$  is  $\|\mathbf{w}\| \cos(\theta) = \frac{1}{\sqrt{5}}$



# Contd...

## Hyperplanes

- In a 3D space, if we have a vector  $\mathbf{w} = [1, 2, 3]^T$  and try to find all points that satisfy  $\mathbf{v} \cdot \mathbf{w} = 1$ , we can obtain a plane that is orthogonal to the vector  $\mathbf{w}$ 
  - The inequalities  $\mathbf{v} \cdot \mathbf{w} > 1$  and  $\mathbf{v} \cdot \mathbf{w} < 1$  again define the two subspaces that are created by the plane



- The same concept applies to high-dimensional spaces as well

# Matrices

## Matrices

- **Matrix** is a rectangular array of real-valued scalars arranged in  $m$  horizontal rows and  $n$  vertical columns
  - Each element  $a_{ij}$  belongs to the  $i^{\text{th}}$  row and  $j^{\text{th}}$  column
  - The elements are denoted  $a_{ij}$  or  $\mathbf{A}_{ij}$  or  $[\mathbf{A}]_{ij}$  or  $\mathbf{A}(i,j)$

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

- For the matrix  $\mathbf{A} \in \mathbb{R}^{m \times n}$ , the size (dimension) is  $m \times n$  or  $(m,n)$ 
  - Matrices are denoted by bold-font upper-case letters

# Contd...

## Matrices

- Addition or subtraction  $(\mathbf{A} \pm \mathbf{B})_{i,j} = \mathbf{A}_{i,j} \pm \mathbf{B}_{i,j}$

$$\begin{bmatrix} 1 & 3 & 1 \\ 1 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 5 \\ 7 & 5 & 0 \end{bmatrix} = \begin{bmatrix} 1+0 & 3+0 & 1+5 \\ 1+7 & 0+5 & 0+0 \end{bmatrix} = \begin{bmatrix} 1 & 3 & 6 \\ 8 & 5 & 0 \end{bmatrix}$$

- Scalar multiplication  $(c\mathbf{A})_{i,j} = c \cdot \mathbf{A}_{i,j}$

$$2 \cdot \begin{bmatrix} 1 & 8 & -3 \\ 4 & -2 & 5 \end{bmatrix} = \begin{bmatrix} 2 \cdot 1 & 2 \cdot 8 & 2 \cdot -3 \\ 2 \cdot 4 & 2 \cdot -2 & 2 \cdot 5 \end{bmatrix} = \begin{bmatrix} 2 & 16 & -6 \\ 8 & -4 & 10 \end{bmatrix}$$

- Matrix multiplication  $(\mathbf{AB})_{i,j} = \mathbf{A}_{i,1}\mathbf{B}_{1,j} + \mathbf{A}_{i,2}\mathbf{B}_{2,j} + \dots + \mathbf{A}_{i,n}\mathbf{B}_{n,j}$

- Defined only if the number of columns of the left matrix is the same as the number of rows of the right matrix
- Note that  $\mathbf{AB} \neq \mathbf{BA}$

$$\begin{bmatrix} \underline{2} & \underline{3} & \underline{4} \\ \underline{1} & \underline{0} & \underline{0} \end{bmatrix} \begin{bmatrix} \underline{0} & \underline{1000} \\ \underline{1} & \underline{100} \\ \underline{0} & \underline{10} \end{bmatrix} = \begin{bmatrix} \underline{3} & \underline{2340} \\ \underline{0} & \underline{1000} \end{bmatrix}$$

# Contd...

## Matrices

- **Transpose** of the matrix:  $\mathbf{A}^T$  has the rows and columns exchanged

$$(\mathbf{A}^T)_{i,j} = \mathbf{A}_{j,i} \qquad \begin{bmatrix} 1 & 2 & 3 \\ 0 & -6 & 7 \end{bmatrix}^T = \begin{bmatrix} 1 & 0 \\ 2 & -6 \\ 3 & 7 \end{bmatrix}$$

- Some properties  $\mathbf{A} + \mathbf{B} = \mathbf{B} + \mathbf{A}$   $\mathbf{A}(\mathbf{B} + \mathbf{C}) = \mathbf{A}\mathbf{B} + \mathbf{A}\mathbf{C}$   
 $(\mathbf{A} + \mathbf{B})^T = \mathbf{A}^T + \mathbf{B}^T$   $\mathbf{A}(\mathbf{B}\mathbf{C}) = (\mathbf{A}\mathbf{B})\mathbf{C}$   
 $(\mathbf{A}^T)^T = \mathbf{A}$   $(\mathbf{A}\mathbf{B})^T = \mathbf{B}^T\mathbf{A}^T$

- **Square matrix**: has the same number of rows and columns
- **Identity matrix** ( $\mathbf{I}_n$ ): has ones on the main diagonal, and zeros elsewhere

- E.g.: identity matrix of size 3×3:  $\mathbf{I}_3 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$



# Contd...

## Matrices

- **Determinant** of a matrix, denoted by  $\det(\mathbf{A})$  or  $|\mathbf{A}|$ , is a real-valued scalar encoding certain properties of the matrix

- E.g., for a matrix of size  $2 \times 2$ : 
$$\det \begin{pmatrix} a & b \\ c & d \end{pmatrix} = ad - bc$$

- For larger-size matrices the determinant of a matrix is calculated as

$$\det(\mathbf{A}) = \sum_{ii} a_{ii} (-1)^{i+i} \det(\mathbf{A}_{ii,ii})$$

- In the above,  $\mathbf{A}_{ii,ii}$  is a **minor** of the matrix obtained by removing the row and column associated with the indices  $i$  and  $j$
- **Trace** of a matrix is the sum of all diagonal elements

$$\text{Tr}(\mathbf{A}) = \sum_{ii} a_{ii}$$

- A matrix for which  $\mathbf{A} = \mathbf{A}^T$  is called a **symmetric matrix**

# Contd...

## Matrices

- Elementwise multiplication of two matrices **A** and **B** is called the *Hadamard product* or *elementwise product*
  - The math notation is  $\odot$

$$\mathbf{A} \odot \mathbf{B} = \begin{bmatrix} a_{11}b_{11} & a_{12}b_{12} & \dots & a_{1n}b_{1n} \\ a_{21}b_{21} & a_{22}b_{22} & \dots & a_{2n}b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1}b_{m1} & a_{m2}b_{m2} & \dots & a_{mn}b_{mn} \end{bmatrix}$$

# Matrix-Vector Products

## Matrices

- Consider a matrix  $\mathbf{A} \in \mathbb{R}^{m \times n}$  and a vector  $\mathbf{x} \in \mathbb{R}^n$
- The matrix can be written in terms of its row vectors (e.g.,  $\mathbf{a}_1^T$  is the first row)

$$\mathbf{A} = \begin{bmatrix} \mathbf{a}_1^T \\ \mathbf{a}_2^T \\ \vdots \\ \mathbf{a}_m^T \end{bmatrix}$$

- The **matrix-vector** product is a column vector of length  $m$ , whose  $i^{\text{th}}$  element is the dot product  $\mathbf{a}_i^T \mathbf{x}$

$$\mathbf{Ax} = \begin{bmatrix} \mathbf{a}_1^T \\ \mathbf{a}_2^T \\ \vdots \\ \mathbf{a}_m^T \end{bmatrix} \mathbf{x} = \begin{bmatrix} \mathbf{a}_1^T \mathbf{x} \\ \mathbf{a}_2^T \mathbf{x} \\ \vdots \\ \mathbf{a}_m^T \mathbf{x} \end{bmatrix}$$

- Note the size:  $\mathbf{A}(m \times n) \mathbf{x}(n \times 1) = \mathbf{Ax}(m \times 1)$

# Contd...

## Matrices

- To multiply two matrices  $\mathbf{A} \in \mathbb{R}^{m \times k}$  and  $\mathbf{B} \in \mathbb{R}^{k \times m}$

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1k} \\ a_{21} & a_{22} & \cdots & a_{2k} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nk} \end{bmatrix}, \quad \mathbf{B} = \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1m} \\ b_{21} & b_{22} & \cdots & b_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ b_{k1} & b_{k2} & \cdots & b_{km} \end{bmatrix}$$

- We can consider the **matrix-matrix product** as dot-products of rows in  $\mathbf{A}$  and columns in  $\mathbf{B}$

$$\mathbf{C} = \mathbf{AB} = \begin{bmatrix} \mathbf{a}_1^\top \\ \mathbf{a}_2^\top \\ \vdots \\ \mathbf{a}_n^\top \end{bmatrix} \begin{bmatrix} \mathbf{b}_1 & \mathbf{b}_2 & \cdots & \mathbf{b}_m \end{bmatrix} = \begin{bmatrix} \mathbf{a}_1^\top \mathbf{b}_1 & \mathbf{a}_1^\top \mathbf{b}_2 & \cdots & \mathbf{a}_1^\top \mathbf{b}_m \\ \mathbf{a}_2^\top \mathbf{b}_1 & \mathbf{a}_2^\top \mathbf{b}_2 & \cdots & \mathbf{a}_2^\top \mathbf{b}_m \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{a}_n^\top \mathbf{b}_1 & \mathbf{a}_n^\top \mathbf{b}_2 & \cdots & \mathbf{a}_n^\top \mathbf{b}_m \end{bmatrix}$$

- Size:  $\mathbf{A}(m \times k) \times \mathbf{B}(k \times m) = \mathbf{C}(m \times m)$

# Inverse of a Matrix

## Matrices

- For a square  $n \times n$  matrix  $\mathbf{A}$  with rank  $n$ ,  $\mathbf{A}^{-1}$  is its *inverse matrix* if their product is an identity matrix  $\mathbf{I}$

$$\mathbf{A}^{-1}\mathbf{A} = \mathbf{A}\mathbf{A}^{-1} = \mathbf{I}$$

- Properties of inverse matrices
$$\left(\mathbf{A}^{-1}\right)^{-1} = \mathbf{A}$$
$$\left(\mathbf{AB}\right)^{-1} = \mathbf{B}^{-1}\mathbf{A}^{-1}$$
- If  $\det(\mathbf{A}) = 0$  (i.e.,  $\text{rank}(\mathbf{A}) < n$ ), then the inverse does not exist
  - A matrix that is not invertible is called a *singular matrix*
- Note that finding an inverse of a large matrix is computationally expensive
  - In addition, it can lead to numerical instability
- If the inverse of a matrix is equal to its transpose, the matrix is said to be *orthogonal matrix*

$$\mathbf{A}^{-1} = \mathbf{A}^T$$

# Pseudo-Inverse of a Matrix

## Matrices

- *Pseudo-inverse* of a matrix
  - Also known as **Moore-Penrose pseudo-inverse**
- For matrices that are not square, the inverse does not exist
  - Therefore, a pseudo-inverse is used
- If  $m > n$ , then the pseudo-inverse is  $\mathbf{A}^\dagger = (\mathbf{A}^\top \mathbf{A})^{-1} \mathbf{A}^\top$  and  $\mathbf{A}^\dagger \mathbf{A} = \mathbf{I}$
- If  $m < n$ , then the pseudo-inverse is  $\mathbf{A}^\dagger = \mathbf{A}^\top (\mathbf{A} \mathbf{A}^\top)^{-1}$  and  $\mathbf{A} \mathbf{A}^\dagger = \mathbf{I}$ 
  - E.g., for a matrix with dimension  $\mathbf{X}_{2 \times 3}$ , a pseudo-inverse can be found of size  $\mathbf{X}_{3 \times 2}^\dagger$ , so that  $\mathbf{X}_{2 \times 3} \mathbf{X}_{3 \times 2}^\dagger = \mathbf{I}_{2 \times 2}$

# Tensors

## *Tensors*

- **Tensors** are  $n$ -dimensional arrays of scalars
  - Vectors are first-order tensors,  $\mathbf{w} \in \mathbb{R}^n$
  - Matrices are second-order tensors,  $\mathbf{A} \in \mathbb{R}^{m \times n}$
  - E.g., a fourth-order tensor is  $\mathbf{T} \in \mathbb{R}^{n_1 \times n_2 \times n_3 \times n_4}$
- Tensors are denoted with upper-case letters of a special font face (e.g.,  $\mathbf{X}$ ,  $\mathbf{Y}$ ,  $\mathbf{Z}$ )
- RGB images are third-order tensors, i.e., as they are 3-dimensional arrays
  - The 3 axes correspond to width, height, and channel
  - E.g.,  $224 \times 224 \times 3$
  - The channel axis corresponds to the color channels (red, green, and blue)

# Manifolds

## *Manifolds*


- Earlier we learned that hyperplanes generalize the concept of planes in high-dimensional spaces
  - Similarly, manifolds can be informally imagined as generalization of the concept of surfaces in high-dimensional spaces
- To begin with an intuitive explanation, the surface of the Earth is an example of a two-dimensional manifold embedded in a three-dimensional space
  - This is true because the Earth looks locally flat, so on a small scale it is like a 2-D plane
  - However, if we keep walking on the Earth in one direction, we will eventually end up back where we started
    - This means that Earth is not really flat, it only looks **locally** like a Euclidean plane, but at large scales it **folds up** on itself, and has a different **global** structure than a flat plane



# Contd...

## *Manifolds*



- Manifolds are studied in mathematics under topological spaces
- An  $n$ -dimensional *manifold* is defined as a topological space with the property that each point has a neighborhood that is homeomorphic to the Euclidean space of dimension  $n$ 
  - This means that a manifold locally resembles Euclidean space near each point
  - Informally, a Euclidean space is locally smooth, it does not have holes, edges, or other sudden changes, and it does not have intersecting neighborhoods
  - Although the manifolds can have very complex structure on a large scale, resemblance of the Euclidean space on a small scale allows to apply standard math concepts
- Examples of 2-dimensional manifolds are shown in the figure
  - The surfaces in the figure have been conveniently cut up into little rectangles that were glued together
  - Those small rectangles locally look like flat Euclidean planes



# Contd...

## *Manifolds*



- Examples of one-dimensional manifolds
  - Upper figure: a circle is a 1-D manifold embedded in 2-D, where each arc of the circle locally resembles a line segment
  - Lower figures: other examples of 1-D manifolds
  - Note that a number 8 figure is not a manifold because it has an intersecting point (it is not Euclidean locally)
- It is hypothesized that in the real-world, high-dimensional data (such as images) lie on low-dimensional manifolds embedded in the high-dimensional space
  - E.g., in ML, let's assume we have a training set of images with size  $224 \times 224 \times 3$  pixels
  - Learning an arbitrary function in such high-dimensional space would be intractable
  - Despite that, all images of the same class ("cats" for example) might lie on a low-dimensional manifold
  - This allows function learning and image classification



# Manifolds

## *Manifolds*

- Example:
  - The data points have 3 dimensions (left figure), i.e., the input space of the data is 3-dimensional
  - The data points lie on a 2-dimensional manifold, shown in the right figure
  - Most ML algorithms extract lower-dimensional data features that enable to distinguish between various classes of high-dimensional input data
    - The low-dimensional representations of the input data are called **embeddings**



# Eigen Decomposition

## *Eigen Decomposition*

- **Eigen decomposition** is decomposing a matrix into a set of eigenvalues and eigenvectors
- **Eigenvalues** of a square matrix  $\mathbf{A}$  are scalars  $\lambda$  and **eigenvectors** are non-zero vectors  $\mathbf{w}$  that satisfy

$$\mathbf{A}\mathbf{w} = \lambda\mathbf{w}$$

- Eigenvalues are found by solving the following equation

$$\det(\mathbf{A} - \lambda\mathbf{I}) = 0$$

- If a matrix  $\mathbf{A}$  has  $n$  linearly independent eigenvectors  $\{\mathbf{w}^1, \dots, \mathbf{w}^n\}$  with corresponding eigenvalues  $\{\lambda_1, \dots, \lambda_n\}$ , the eigen decomposition of  $\mathbf{A}$  is given by

$$\mathbf{A} = \mathbf{W}\mathbf{\Lambda}\mathbf{W}^{-1}$$

- Columns of the matrix  $\mathbf{W}$  are the eigenvectors, i.e.,  $\mathbf{W} = [\mathbf{w}^1, \dots, \mathbf{w}^n]$
  - $\mathbf{\Lambda}$  is a diagonal matrix of the eigenvalues, i.e.,  $\mathbf{\Lambda} = [\lambda_1, \dots, \lambda_n]$
- To find the inverse of the matrix  $\mathbf{A}$ , we can use  $\mathbf{A}^{-1} = \mathbf{W}\mathbf{\Lambda}^{-1}\mathbf{W}^{-1}$ 
  - This involves simply finding the inverse  $\mathbf{\Lambda}^{-1}$  of a diagonal matrix

# Contd...

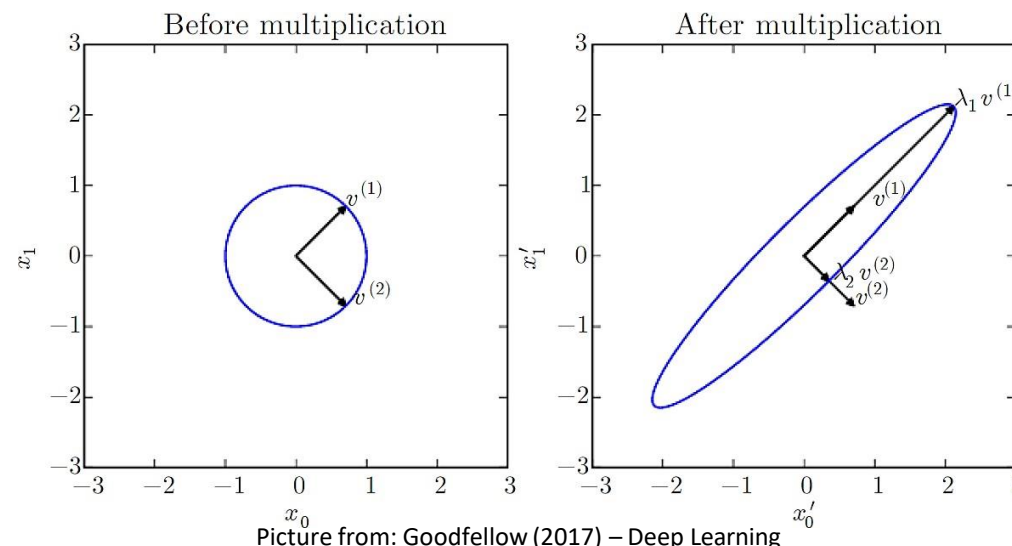
## *Eigen Decomposition*

- Decomposing a matrix into eigenvalues and eigenvectors allows to analyze certain properties of the matrix
  - If all eigenvalues are positive, the matrix is **positive definite**
  - If all eigenvalues are positive or zero-valued, the matrix is **positive semidefinite**
  - If all eigenvalues are negative or zero-values, the matrix is **negative semidefinite**
    - Positive semidefinite matrices are interesting because they guarantee that  $\forall \mathbf{x}, \mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$
- Eigen decomposition can also simplify many linear-algebraic computations
  - The determinant of A can be calculated as
$$\det(\mathbf{A}) = \lambda_1 \lambda_2 \cdots \lambda_n$$
    - If any of the eigenvalues are zero, the matrix is singular (it does not have an inverse)
- However, not every matrix can be decomposed into eigenvalues and eigenvectors
  - Also, in some cases the decomposition may involve complex numbers
  - Still, every real symmetric matrix is guaranteed to have an eigen decomposition according to  $\mathbf{A} = \mathbf{W} \mathbf{\Lambda} \mathbf{W}^{-1}$ , where  $\mathbf{W}$  is an orthogonal matrix

# Contd...

## *Eigen Decomposition*

- Geometric interpretation of the eigenvalues and eigenvectors is that they allow to stretch the space in specific directions
  - Left figure: the two eigenvectors  $\mathbf{w}^1$  and  $\mathbf{w}^2$  are shown for a matrix, where the two vectors are unit vectors (i.e., they have a length of 1)
  - Right figure: the vectors  $\mathbf{w}^1$  and  $\mathbf{w}^2$  are multiplied with the eigenvalues  $\lambda_1$  and  $\lambda_2$ 
    - We can see how the space is scaled in the direction of the larger eigenvalue  $\lambda_1$
- E.g., this is used for dimensionality reduction with PCA (principal component analysis) where the eigenvectors corresponding to the largest eigenvalues are used for extracting the most important data dimensions





# Differential Calculus

## Differential Calculus

- For a function  $f: \mathbb{R} \rightarrow \mathbb{R}$ , the **derivative** of  $f$  is defined as

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- If  $f'(a)$  exists,  $f$  is said to be **differentiable** at  $a$
- If  $f'$  is differentiable for  $\forall x \in [a, b]$ , then  $f$  is differentiable on this interval
  - We can also interpret the derivative  $f'(x)$  as the **instantaneous rate of change** of  $f(x)$  with respect to  $x$
  - I.e., for a small change in  $x$ , what is the rate of change of  $f(x)$
- Given  $y = f(x)$ , where  $x$  is an independent variable and  $y$  is a dependent variable, the following expressions are equivalent:

$$f'(x) = f' = \frac{dy}{dx} = \frac{df}{dx} = \frac{d}{dx} f(x) = Df(x) = D_x f(x)$$

- The symbols  $\frac{d}{dx}$ ,  $D$ , and  $D_x$  are **differentiation operators** that indicate operation of **differentiation**

# Contd...

## *Differential Calculus*

- The following rules are used for computing the derivatives of explicit functions
  - **Derivative of constants.**  $\frac{d}{dx}c = 0$ .
  - **Derivative of linear functions.**  $\frac{d}{dx}(ax) = a$ .
  - **Power rule.**  $\frac{d}{dx}x^n = nx^{n-1}$ .
  - **Derivative of exponentials.**  $\frac{d}{dx}e^x = e^x$ .
  - **Derivative of the logarithm.**  $\frac{d}{dx}\log(x) = \frac{1}{x}$ .
  - **Sum rule.**  $\frac{d}{dx}(g(x) + h(x)) = \frac{dg}{dx}(x) + \frac{dh}{dx}(x)$ .
  - **Product rule.**  $\frac{d}{dx}(g(x) \cdot h(x)) = g(x)\frac{dh}{dx}(x) + \frac{dg}{dx}(x)h(x)$ .
  - **Chain rule.**  $\frac{d}{dx}g(h(x)) = \frac{dg}{dh}(h(x)) \cdot \frac{dh}{dx}(x)$ .



# Contd...

## Differential Calculus

- The derivative of the first derivative of a function  $ff(x)$  is the *second derivative* of  $ff(x)$

$$\frac{d^2ff}{dx^2} = \frac{d}{dx} \left( \frac{dff}{dx} \right)$$

- The second derivative quantifies how the rate of change of  $ff(x)$  is changing
  - E.g., in physics, if the function describes the displacement of an object, the first derivative gives the velocity of the object (i.e., the rate of change of the position)
  - The second derivative gives the acceleration of the object (i.e., the rate of change of the velocity)
- If we apply the differentiation operation any number of times, we obtain the *n-th derivative* of  $ff(x)$

$$ff^{(m)}(x) = \frac{d^mff}{dx^m} = \left( \frac{d}{dx} \right)^m ff(x)$$


# Geometric Interpretation

## Differential Calculus

- To provide a geometric interpretation of the derivatives, let's consider a first-order Taylor series approximation of  $f(x)$  at  $x = x_0$

$$f(x) \approx f(x_0) + \frac{df}{dx}\bigg|_{x_0} (x - x_0)$$

- The expression approximates the function  $f(x)$  by a line which passes through the point  $(x_0, f(x_0))$  and has slope  $\frac{df}{dx}\bigg|_{x_0}$  (i.e., the value of  $\frac{df}{dx}$  at the point  $x_0$ )
- Therefore, the first derivative of a function is also the **slope of the tangent line** to the curve of the function



# Partial Derivatives

## Differential Calculus

- So far, we looked at functions of a single variable, where  $f: \mathbb{R} \rightarrow \mathbb{R}$
- Functions that depend on many variables are called **multivariate functions**
- Let  $y = f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$  be a multivariate function with  $n$  variables
  - The input is an  $n$ -dimensional vector  $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$  and the output is a scalar  $y$
  - The mapping is  $f: \mathbb{R}^n \rightarrow \mathbb{R}$
- The **partial derivative** of  $y$  with respect to its  $i^{\text{th}}$  parameter  $x_i$  is

$$\frac{\partial y}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(x_1, x_2, \dots, x_i + h, \dots, x_n) - f(x_1, x_2, \dots, x_i, \dots, x_n)}{h}$$

- To calculate  $\frac{\partial}{\partial x_i}$  ( $\partial$  pronounced “del” or we can just say “partial derivative”), we can treat  $x_1, x_2, \dots, x_{i-1}, x_{i+1}, \dots, x_n$  as constants and calculate the derivative of  $y$  only with respect to  $x_i$
- For notation of partial derivatives, the following are equivalent:

$$\frac{\partial y}{\partial x_i} = \frac{\partial f}{\partial x_i} = \frac{\partial}{\partial x_i} f(\mathbf{x}) = f_{x_i} = f_i = D_i f = D_{x_i} f$$

# Gradient

## Differential Calculus

- We can concatenate partial derivatives of a multivariate function with respect to all its input variables to obtain the **gradient** vector of the function
- The gradient of the multivariate function  $f(\mathbf{x})$  with respect to the  $n$ -dimensional input vector  $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$ , is a vector of  $n$  partial derivatives

$$\nabla f(\mathbf{x}) = \left[ \frac{\partial f(\mathbf{x})}{\partial x_1}, \frac{\partial f(\mathbf{x})}{\partial x_2}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_n} \right]^T$$

- When there is no ambiguity, the notations  $\nabla f(\mathbf{x})$  or  $\nabla f$  are often used for the gradient instead of  $\nabla f(\mathbf{x})$ 
  - The symbol for the gradient is the Greek letter  $\nabla$  (pronounced “nabla”), although  $\nabla f(\mathbf{x})$  is more often it is pronounced “gradient of  $f$  with respect to  $\mathbf{x}$ ”
- In ML, the gradient descent algorithm relies on the opposite direction of the gradient of the loss function  $\mathcal{L}$  with respect to the model parameters  $\theta$  ( $\nabla_{\theta} \mathcal{L}$ ) for minimizing the loss function
  - Adversarial examples can be created by adding perturbation in the direction of the gradient of the loss  $\mathcal{L}$  with respect to input examples  $\mathbf{x}$  ( $\nabla_{\mathbf{x}} \mathcal{L}$ ) for maximizing the loss function

# Optimization


## *Optimization*

- **Optimization** is concerned with optimizing an **objective function** — finding the value of an argument that minimizes or maximizes the function
  - Most optimization algorithms are formulated in terms of minimizing a function  $f(x)$
  - Maximization is accomplished via minimizing the negative of an objective function (e.g., minimize  $-f(x)$ )
  - In minimization problems, the objective function is often referred to as a **cost function** or **loss function** or **error function**
- Optimization is very important for machine learning
  - The performance of optimization algorithms affect the model's training efficiency
- Most optimization problems in machine learning are **nonconvex**
  - Meaning that the loss function is not a convex function
  - Nonetheless, the design and analysis of algorithms for solving convex problems has been very instructive for advancing the field of machine learning

# Contd...

## Optimization


- Optimization and machine learning have related, but somewhat different goals
  - Goal in optimization: minimize an objective function
    - For a set of training examples, reduce the **training error**
  - Goal in ML: find a suitable model, to predict on data examples
    - For a set of testing examples, reduce the **generalization error**
- For a given empirical function  $g$  (dashed purple curve), optimization algorithms attempt to find the point of minimum **empirical risk**
- The expected function  $f$  (blue curve) is obtained given a limited amount of training data examples
- ML algorithms attempt to find the point of minimum **expected risk**, based on minimizing the error on a set of testing examples
  - Which may be at a different location than the minimum of the training examples
  - And which may not be minimal in a formal sense



# Stationary Points

## Optimization

- **Stationary points** ( or **critical points**) of a differentiable function  $f(x)$  of one variable are the points where the derivative of the function is zero, i.e.,  $f'(x) = 0$
- The stationary points can be:
  - **Minimum**, a point where the derivative changes from negative to positive
  - **Maximum**, a point where the derivative changes from positive to negative
  - **Saddle point**, derivative is either positive or negative on both sides of the point
- The minimum and maximum points are collectively known as **extremum points**
- The nature of stationary points can be determined based on the second derivative of  $f(x)$  at the point
  - If  $f''(x) > 0$ , the point is a minimum
  - If  $f''(x) < 0$ , the point is a maximum
  - If  $f''(x) = 0$ , inconclusive, the point can be a saddle point, but it may not
- The same concept also applies to gradients of multivariate functions






# Local Minima

## Optimization

- Among the challenges in optimization of model's parameters in ML involve local minima, saddle points, vanishing gradients
- For an objective function  $f(x)$ , if the value at a point  $x$  is the minimum of the objective function **over the entire domain** of  $x$ , then it is the *global minimum*
- If the value of  $f(x)$  at  $x$  is smaller than the values of the objective function at any other points in **the vicinity** of  $x$ , then it is the *local minimum*
- The objective functions in ML usually have many local minima
  - When the solution of the optimization algorithm is near the local minimum, the gradient of the loss function approaches or becomes zero (vanishing gradients)
  - Therefore, the obtained solution in the final iteration can be a local minimum, rather than the global minimum







# Saddle Points

## Optimization


- The gradient of a function  $f(x)$  at a **saddle point** is 0, but the point is not a minimum or maximum point
  - The optimization algorithms may stall at saddle points, without reaching a minima
- Note also that the point of a function at which the sign of the curvature changes is called an **inflection point**
  - An inflection point ( $f''(x) = 0$ ) can also be a saddle point, but it does not have to be
- For the 2D function (right figure), the saddle point is at (0,0)
  - The point looks like a saddle, and gives the minimum with respect to  $x$ , and the maximum with respect to  $y$




# Convex Optimization

## Optimization

- A function of a single variable is *concave* if every line segment joining two points on its graph does not lie above the graph at any point
- Symmetrically, a function of a single variable is *convex* if every line segment joining two points on its graph does not lie below the graph at any point



A concave function:  
no line segment joining  
two points on the graph  
lies above the graph  
at any point



A convex function:  
no line segment joining  
two points on the graph  
lies below the graph  
at any point




A function that is neither  
concave nor convex:  
the line segment shown lies  
above the graph at some  
points and below it at others

# Contd...

## Optimization

- In mathematical terms, the function  $f$  is a **convex function** if for all points  $x_1, x_2$  and for all  $\lambda \in [0, 1]$

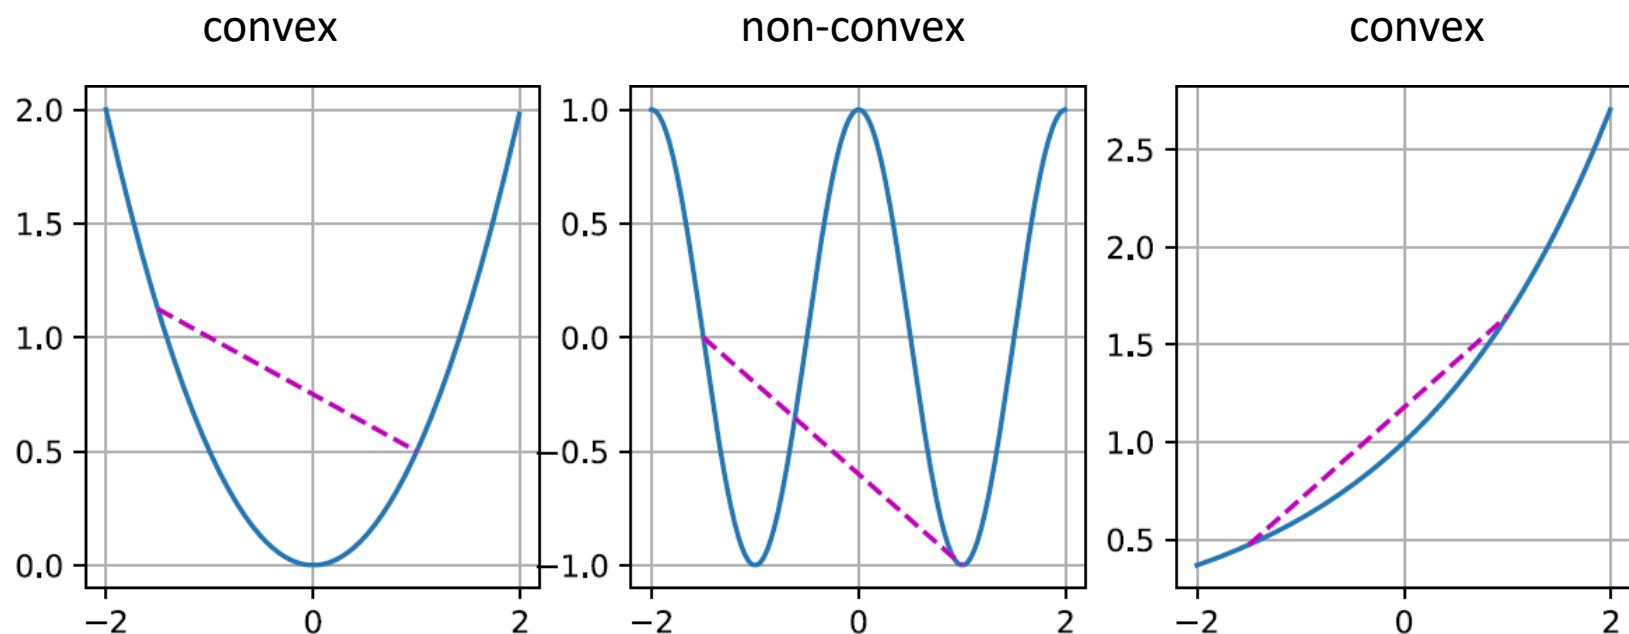
$$\lambda f(x_1) + (1 - \lambda)f(x_2) \geq f(\lambda x_1 + (1 - \lambda)x_2)$$



# Convex Functions

## Optimization

- One important property of convex functions is that they **do not have local minima**
  - Every local minimum of a convex function is a global minimum
  - I.e., every point at which the gradient of a convex function = 0 is the global minimum
  - The figure below illustrates two convex functions, and one nonconvex function



# Probability

## *Probability*

- Intuition:
  - In a process, several outcomes are possible
  - When the process is repeated a large number of times, each outcome occurs with a *relative frequency*, or *probability*
  - If a particular outcome occurs more often, we say it is more probable
- Probability arises in two contexts
  - In actual repeated experiments
    - Example: You record the color of 1,000 cars driving by. 57 of them are green. You **estimate** the probability of a car being green as  $57/1,000 = 0.057$ .
  - In idealized conceptions of a repeated process
    - Example: You consider the behavior of an unbiased six-sided die. The **expected** probability of rolling a 5 is  $1/6 = 0.1667$ .
    - Example: You need a model for how people's heights are distributed. You choose a normal distribution to represent the **expected** relative probabilities.

# Contd...

## *Probability*

- Solving machine learning problems requires to deal with uncertain quantities, as well as with stochastic (non-deterministic) quantities
  - Probability theory provides a mathematical framework for representing and quantifying uncertain quantities
- There are different sources of uncertainty:
  - Inherent stochasticity in the system being modeled
    - For example, most interpretations of quantum mechanics describe the dynamics of subatomic particles as being probabilistic
  - Incomplete observability
    - Even deterministic systems can appear stochastic when we cannot observe all of the variables that drive the behavior of the system
  - Incomplete modeling
    - When we use a model that must discard some of the information we have observed, the discarded information results in uncertainty in the model's predictions
    - E.g., discretization of real-numbered values, dimensionality reduction, etc.

# Random variables

## Probability

- A **random variable**  $X$  is a variable that can take on different values
  - Example:  $X$  = rolling a die
    - Possible values of  $X$  comprise the **sample space**, or **outcome space**,  $\mathcal{S} = \{1, 2, 3, 4, 5, 6\}$
    - We denote the event of “seeing a 5” as  $\{X = 5\}$  or  $X = 5$
    - The probability of the event is  $P(\{X = 5\})$  or  $P(X = 5)$
    - Also,  $P(5)$  can be used to denote the probability that  $X$  takes the value of 5
- A **probability distribution** is a description of how likely a random variable is to take on each of its possible states
  - A compact notation is common, where  $P(X)$  is the probability distribution over the random variable  $X$ 
    - Also, the notation  $X \sim P(X)$  can be used to denote that the random variable  $X$  has probability distribution  $P(X)$
- Random variables can be discrete or continuous
  - **Discrete random variables** have finite number of states: e.g., the sides of a die
  - **Continuous random variables** have infinite number of states: e.g., the height of a person



# Axioms of probability

## *Probability*

- The probability of an event  $\mathcal{A}$  in the given sample space  $\mathcal{S}$ , denoted as  $P(\mathcal{A})$ , must satisfy the following properties:
  - Non-negativity
    - For any event  $\mathcal{A} \in \mathcal{S}$ ,  $P(\mathcal{A}) \geq 0$
  - All possible outcomes
    - Probability of the entire sample space is 1,  $P(\mathcal{S}) = 1$
  - Additivity of disjoint events
    - For all events  $\mathcal{A}_1, \mathcal{A}_2 \in \mathcal{S}$  that are mutually exclusive ( $\mathcal{A}_1 \cap \mathcal{A}_2 = \emptyset$ ), the probability that both events happen is equal to the sum of their individual probabilities,  $P(\mathcal{A}_1 \cup \mathcal{A}_2) = P(\mathcal{A}_1) + P(\mathcal{A}_2)$
- The probability of a random variable  $P(X)$  must obey the axioms of probability over the possible values in the sample space  $\mathcal{S}$





# Discrete Variables

## Probability

- A probability distribution over **discrete variables** may be described using a *probability mass function* (PMF)
  - E.g., sum of two dice
- A probability distribution over **continuous variables** may be described using a *probability density function* (PDF)
  - E.g., waiting time between eruptions of Old Faithful
  - A PDF gives the probability of an infinitesimal region with volume  $dX$
  - To find the probability over an interval  $[a, b]$ , we can integrate the PDF as follows:

$$P(X \in [a, b]) = \int_a^b P(X) dX$$



# Multivariate Random Variables


## Probability

- We may need to consider several random variables at a time
  - If several random processes occur in parallel or in sequence
  - E.g., to model the relationship between several diseases and symptoms
  - E.g., to process images with millions of pixels (each pixel is one random variable)
- Next, we will study probability distributions defined over multiple random variables
  - These include joint, conditional, and marginal probability distributions
- The individual random variables can also be grouped together into a random vector, because they represent different properties of an individual statistical unit
- A *multivariate random variable* is a vector of multiple random variables  $\mathbf{X} = (X_1, X_2, \dots, X_m)^T$

# Joint Probability Distribution

## Probability

- Probability distribution that acts on many variables at the same time is known as a *joint probability distribution*
- Given any values  $x$  and  $y$  of two random variables  $X$  and  $Y$ , what is the probability that  $X = x$  and  $Y = y$  simultaneously?
  - $PP(X = x, Y = y)$  denotes the joint probability
  - We may also write  $PP(x, y)$  for brevity




# Marginal Probability Distribution

## Probability

- **Marginal probability distribution** is the probability distribution of a single variable
  - It is calculated based on the joint probability distribution  $P(X,Y)$
  - I.e., using the **sum rule**:  $P(X=x) = \sum_y P(X=x, Y=y)$ 
    - For continuous random variables, the summation is replaced with integration,  $P(X=x) = \int P(X=x, Y=y) dy$
  - This process is called **marginalization**

marginal probability:  $p(X = \text{minivan}) = 0.0741 + 0.1111 + 0.1481 = 0.3333$




# Conditional Probability Distribution

## Probability

- **Conditional probability distribution** is the probability distribution of one variable provided that another variable has taken a certain value
  - Denoted  $P(X = x | Y = y)$
- Note that:  $P(X = x | Y = y) = \frac{P(X=x, Y=y)}{P(Y=y)}$

conditional probability:  $p(Y = \text{European} | X = \text{minivan}) = 0.1481 / (0.0741 + 0.1111 + 0.1481) = 0.4433$



# Bayes' Theorem

## Probability

- **Bayes' theorem** – allows to calculate conditional probabilities for one variable when conditional probabilities for another variable are known

$$P(X|Y) = \frac{P(Y|X)P(X)}{P(Y)}$$

- Also known as Bayes' rule
- **Multiplication rule** for the joint distribution is used:  $P(X,Y) = P(Y|X)P(X)$
- By symmetry, we also have:  $P(Y,X) = P(X|Y)P(Y)$
- The terms are referred to as:
  - $P(X)$ , the **prior probability**, the initial degree of belief for  $X$
  - $P(X|Y)$ , the **posterior probability**, the degree of belief after incorporating the knowledge of  $Y$
  - $P(Y|X)$ , the **likelihood** of  $Y$  given  $X$
  - $P(Y)$ , the **evidence**
  - Bayes' theorem:  $\text{posterior probability} = \frac{\text{likelihood} \times \text{prior probability}}{\text{evidence}}$



# Independence

## Probability

- Two random variables  $X$  and  $Y$  are *independent* if the occurrence of  $Y$  does not reveal any information about the occurrence of  $X$ 
  - E.g., two successive rolls of a die are independent
- Therefore, we can write:  $P(X|Y) = P(X)$ 
  - The following notation is used:  $X \perp Y$
  - Also note that for independent random variables:  $P(X,Y) = P(X)P(Y)$
- In all other cases, the random variables are *dependent*
  - E.g., duration of successive eruptions of Old Faithful
  - Getting a king on successive draws from a deck (the drawn card is not replaced)
- Two random variables  $X$  and  $Y$  are *conditionally independent* given another random variable  $Z$  if and only if  $P(X,Y|Z) = P(X|Z)P(Y|Z)$ 
  - This is denoted as  $X \perp Y | Z$

# Expected Value

## Probability

- The *expected value* or *expectation* of a function  $f(X)$  with respect to a probability distribution  $P(X)$  is the average (mean) when  $X$  is drawn from  $P(X)$
- For a discrete random variable  $X$ , it is calculated as

$$\mathbb{E}_{X \sim P}[f(X)] = \sum_X P(X) f(X)$$

- For a continuous random variable  $X$ , it is calculated as

$$\mathbb{E}_{X \sim P}[f(X)] = \int P(X) f(X) dX$$

- When the identity of the distribution is clear from the context, we can write  $\mathbb{E}_X[f(X)]$
- If it is clear which random variable is used, we can write just  $\mathbb{E}[f(X)]$
- Mean is the most common measure of central tendency of a distribution
  - For a random variable:  $f(X_i) = X_i \Rightarrow \mu = \mathbb{E}[X_i] = \sum_i P(X_i) X_i$
  - This is similar to the mean of a sample of observations:  $\mu = \frac{1}{N} \sum_i X_i$
  - Other measures of central tendency: median, mode



# Variance

## Probability

- **Variance** gives the measure of how much the values of the function  $ff(X)$  deviate from the expected value as we sample values of  $X$  from  $P(X)$

$$\text{Var}(ff(X)) = \mathbb{E}[(ff(X) - \mathbb{E}[ff(X)])^2]$$

- When the variance is low, the values of  $ff(X)$  cluster near the expected value
- Variance is commonly denoted with  $\sigma^2$ 
  - The above equation is similar to a function  $ff(X_i) = X_i - \mu$
  - We have  $\sigma^2 = \sum_{ii} P(X_i) (X_i - \mu)^2$
  - This is similar to the formula for calculating the variance of a sample of observations:  
$$\sigma^2 = \frac{1}{N-1} \sum_{ii} (X_i - \mu)^2$$
- The square root of the variance is the **standard deviation**
  - Denoted  $\sigma = \sqrt{\text{Var}(X)}$




# Covariance

## Probability

- **Covariance** gives the measure of how much two random variables are linearly related to each other

$$\text{Cov}(f(X), g(Y)) = \mathbb{E}[(f(X) - \mathbb{E}[f(X)])(g(Y) - \mathbb{E}[g(Y)])]$$

- If  $f(X_i) = X_i - \mu_X$  and  $g(Y_i) = Y_i - \mu_Y$ 
  - Then, the covariance is:  $\text{Cov}(X, Y) = \sum_{ii} P(X_i, Y_i) (X_i - \mu_X) (Y_i - \mu_Y)$
  - Compare to covariance of actual samples:  $\text{Cov}(X, Y) = \frac{1}{N-1} \sum_{ii} (X_i - \mu_X) (Y_i - \mu_Y)$
- The covariance measures the tendency for  $X$  and  $Y$  to deviate from their means in same (or opposite) directions at same time




# Correlation

## Probability

- **Correlation coefficient** is the covariance normalized by the standard deviations of the two variables

$$\text{corr}(X,Y) = \frac{\text{Cov}(X,Y)}{\sigma_X \sigma_Y}$$

- It is also called **Pearson's correlation coefficient** and it is denoted  $\rho(X,Y)$
- The values are in the interval  $[-1, 1]$
- It only reflects linear dependence between variables, and it does not measure non-linear dependencies between the variables



Linear dependence  
with noise

Linear dependence  
without noise


Various nonlinear  
dependencies

# Probability Distributions

## Probability


- *Bernoulli distribution*

- Binary random variable  $X$  with states  $\{0, 1\}$
- The random variable can encode a coin flip which comes up 1 with probability  $p$  and 0 with probability  $1 - p$
- Notation:  $X \sim \text{Bernoulli}(p)$



- *Uniform distribution*

- The probability of each value  $i \in \{1, 2, \dots, m\}$  is  $\frac{1}{m}$
- Notation:  $X \sim U(m)$
- Figure:  $m = 5$ ,  $p = 0.2$



# Contd...



## Probability

### • Binomial distribution

- Performing a sequence of  $n$  independent experiments, each of which has probability  $p$  of succeeding, where  $p \in \{0, 1\}$
- The probability of getting  $k$  successes in  $n$  trials is  $P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}$
- Notation:  $X \sim \text{Bin}(n, p)$

### • Poisson distribution

- A number of events occurring independently in a fixed interval of time with a known rate  $\lambda$
- A discrete random variable  $X$  with states  $k \in \{0, 1, 2, \dots\}$  has probability  $P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$
- The rate  $\lambda$  is the average number of occurrences of the event
- Notation:  $X \sim \text{Poi}(\lambda)$





# Contd...

## Probability

- *Gaussian distribution*

- The most well-studied distribution
  - Referred to as **normal distribution** or informally **bell-shaped distribution**
- Defined with the mean  $\mu$  and variance  $\sigma^2$
- Notation:  $X \sim \mathcal{N}(\mu, \sigma^2)$
- For a random variable  $X$  with  $n$  independent measurements, the density is

$$P_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



# Cross-entropy

## Information Theory

- **Cross-entropy** is closely related to the KL divergence, and it is defined as the summation of the entropy  $H(P)$  and KL divergence  $D_{KL}(P||Q)$

$$\mathcal{C}(P, Q) = H(P) + D_{KL}(P||Q)$$

- Alternatively, the cross-entropy can be written as

$$\mathcal{C}(P, Q) = -\mathbb{E}_{X \sim P} [\log Q(X)]$$

- In machine learning, let's assume a classification problem based on a set of data examples  $\{x_1, x_2, \dots, x_n\}$ , that need to be classified into  $k$  classes
  - For each data example  $x_i$  we have a class label  $y_i$ 
    - The true labels  $y$  follow the true distribution  $P$
  - The goal is to train a classifier (e.g., a NN) parameterized by  $\theta$ , that outputs a predicted class label  $\hat{y}_i$  for each data example  $x_i$ 
    - The predicted labels  $\hat{y}$  follow the estimated distribution  $Q$
  - The cross-entropy loss between the true distribution  $P$  and the estimated distribution  $Q$  is calculated as:  $\mathcal{C}(y, \hat{y}) = -\mathbb{E}_{X \sim P} [\log Q(X)] = -\sum_X P(X) \log Q(X) = -\sum_i y_i \log \hat{y}_i$ 
    - The further away the true and estimated distributions are, the greater the cross-entropy loss is



